

All the code can be accessed here and executed in google colab or any Jupyter Notebook environment:  
<https://github.com/DragonBoy25830/caltech-cs-156>

1)

```

Ein_threshold = 0.008
sigma = 0.1
d = 8
[10] ✓ 0.0s

def lin_reg_error(sigma_in, d_in, N):
    return sigma_in * sigma_in * (1 - (d_in + 1) / N)
[11] ✓ 0.0s

for i in range(1, 1000):
    if lin_reg_error(sigma, d, i) > Ein_threshold:
        print(f"Minimum N: {i}")
        break
[12] ✓ 0.0s

... Minimum N: 45


```

The above code shows us that the smallest  $N$  that leads to the expected in-sample error being greater than 0.008 is 45.

The next answer choice larger than this is 100, so

the answer is c

2)



In a non-linear transformation:

$$g(x) = \text{sign}(\tilde{w} \Phi(x))$$

$$= \text{sign}\left([\tilde{w}_0 \ \tilde{w}_1 \ \tilde{w}_2] \cdot [1 \ x_1^2 \ x_2^2]\right)$$

$$= \text{sign}(\tilde{w}_0 + \tilde{w}_1 x_1^2 + \tilde{w}_2 x_2^2)$$

If we fix  $x_2=0$ , we see that  $g(x) > 0$  when  $x_1$  is large and  $g(x) < 0$  when  $x_1$  is small. Assuming that  $\tilde{w}_0$  is made sufficiently large to achieve the desired boundary,  $\tilde{w}_1 < 0$  to produce the desired decision boundary.

If we fix  $x_1=0$ , we see that  $g(x) > 0$  for all  $x_2$ . As such  $\tilde{w}_2 > 0$ .

Since  $\tilde{w}_1 < 0$  &  $\tilde{w}_2 > 0$ , the answer is d.

3) In the transformed space,  $d_{vc} \leq \tilde{d} + 1$  where  $\tilde{d}$  is the number of features in the transformed space. By simple counting, we see that  $\Phi_4: x$  has 15 features, so  $d_{vc} < 14 + 1 = 15$ . The smallest value among the choices not smaller than the VC dimension of a linear model in the transformed space is 15, so the answer is c.

4) 
$$\frac{\partial E}{\partial u} [(ue^v - 2ve^{-u})^2] = 2(ue^v - 2ve^{-u})(e^v + 2ve^{-u})$$

The answer is e.

5)

```
import numpy as np

[1] ✓ 0.1s

n = 0.1
uv_array = [(1, 1)]
[2] ✓ 0.0s

def partial_Eu(u, v):
    return 2 * (u * np.exp(v) - 2 * v * np.exp(-1 * u)) * (np.exp(v) + 2 * v * np.exp(-1 * u))
[3] ✓ 0.0s

def partial_Ev(u, v):
    return 2 * (u * np.exp(v) - 2 * v * np.exp(-1 * u)) * (u * np.exp(v) - 2 * np.exp(-1 * u))
[4] ✓ 0.0s

def get_last_uv():
    return uv_array[-1][0], uv_array[-1][1]
[5] ✓ 0.0s

def gradient_descent(n_in):
    u_t, v_t = get_last_uv()

    u_t1 = u_t - n_in * partial_Eu(u_t, v_t)
    v_t1 = v_t - n_in * partial_Ev(u_t, v_t)

    uv_array.append((u_t1, v_t1))
[6] ✓ 0.0s

def calc_error(u, v):
    return (u * np.exp(v) - 2 * v * np.exp(-1 * u)) ** 2
[7] ✓ 0.0s

num_iter = 0

while True:
    u_t, v_t = get_last_uv()

    if calc_error(u_t, v_t) < 1e-14:
        print(num_iter)
        break

    gradient_descent(n)
    num_iter += 1
[8] ✓ 0.0s

... 10
```

The above code shows that it takes 10 iterations for

the error  $E(u,v)$  to fall below  $10^{-14}$ , so the answer is d

6)

```
get_last_uv()
[9] ✓ 0.0s
... (0.04473629039778207, 0.023958714099141746)
```

The above code outputs the final  $(u,v)$  to be closest to  $(0.045, 0.024)$ , so the answer is e

7)

```
uv_array = [(1, 1)]
[10] ✓ 0.0s

def coordinate_descent(n_in):
    u_t, v_t = get_last_uv()

    u_t1 = u_t - n_in * partial_Eu(u_t, v_t)
    v_t1 = v_t - n_in * partial_Ev(u_t1, v_t)

    uv_array.append((u_t1, v_t1))
[11] ✓ 0.0s

for i in range(15):
    coordinate_descent(n)

    if i == 14:
        u_t, v_t = get_last_uv()
        print(calc_error(u_t, v_t))
[12] ✓ 0.0s
... 0.13981379199615315
```

The error outputted by the code above is closest to  $10^{-1}$ , so the answer is a

8)

```
import numpy as np
import random as rand
[1] ✓ 0.0s

num_train_points = 100
num_test_points = 1000
learning_rate = 0.01
[2] ✓ 0.0s
```

8)

```

import numpy as np
import random as rand

[1] ✓ 0.0s

>
num_train_points = 100
num_test_points = 1000
learning_rate = 0.01

[2] ✓ 0.0s

def generate_target_function():
    x0 = rand.uniform(-1, 1)
    y0 = rand.uniform(-1, 1)
    x1 = rand.uniform(-1, 1)
    y1 = rand.uniform(-1, 1)

    m_f = (y1 - y0) / (x1 - x0)
    b_f = y0 - m_f * x0
    f = [m_f, b_f]

    return f

[3] ✓ 0.0s

def classify_point(m, b, x, y):
    expected_value = m * x + b

    if y >= expected_value:
        return 1
    else:
        return -1

[4] ✓ 0.0s

```

```

def generate_data(m, b, num_points):
    bias = [1 for _ in range(num_points)]
    x1 = [rand.uniform(-1, 1) for _ in range(num_points)]
    x2 = [rand.uniform(-1, 1) for _ in range(num_points)]

    xn_pre_bias = np.column_stack((x1, x2))
    xn = np.column_stack((bias, xn_pre_bias))
    yn = [classify_point(m, b, x1, x2) for (bias, x1, x2) in xn]

    training_data = np.column_stack((xn, yn))

    return training_data

[5] ✓ 0.0s

def norm(vec):
    return np.sqrt(np.sum(vec * vec))

[6] ✓ 0.0s

>
def SGD(input_weight, training_data):
    perm_training_data = np.random.permutation(training_data)
    xn = np.transpose(np.delete(np.transpose(perm_training_data), 3, 0))
    yn = np.transpose(np.transpose(perm_training_data)[3])

    N = len(xn)

    for i in range(N):
        xi = xn[i]
        yi = yn[i]

        gradient = -1 * (yi * xi) / (1 + np.exp(yi * np.dot(input_weight, xi)))
        input_weight = input_weight - learning_rate * gradient

    return input_weight

[7] ✓ 0.0s

```

```

def calc_cross_entropy_error(testing_data, weight):
    xn = np.transpose(np.delete(np.transpose(testing_data), 3, 0))
    yn = np.transpose(np.transpose(testing_data)[3])

    N = len(xn)

    error = 0

    for i in range(N):
        xi = xn[i]
        yi = yn[i]

        error += np.log(1 + np.exp(-yi * np.dot(weight, xi)))

    return (1 / N) * error

[8] ✓ 0.0s

Eout_errors = []
epoch_weights = np.array([[0, 0, 0]])
num_epochs_arr = []

[9] ✓ 0.0s

```

```

def run_experiment():
    m_f, b_f = generate_target_function()
    # generate target function
    train_data = generate_data(m_f, b_f, num_train_points)

    # run SGD
    # run N epochs till subsequent weights are within two decimal places

    prev_weight = epoch_weights[-1]
    weight = SGD(prev_weight, train_data)
    epoch_iter = 1

    while norm(prev_weight - weight) >= 0.01:
        np.append(epoch_weights, [weight], axis=0)
        prev_weight = weight
        weight = SGD(prev_weight, train_data)
        epoch_iter += 1

    num_epochs_arr.append(epoch_iter)
    # generate testing data
    testing_data = generate_data(m_f, b_f, num_test_points)
    # estimate E_out
    Eout_errors.append(calc_cross_entropy_error(testing_data, weight))

[12] ✓ 0.0s

    for i in range(100):
        run_experiment()

[13] ✓ 15.8s

    np.mean(Eout_errors)

[14] ✓ 0.0s

... 0.10069016921792245

```

The above code outputs a mean  $E_{out}$  for  $N=100$  of 0.10069 which is closest to 0.100, so the answer is d

9)

```

    np.mean(num_epochs_arr)

[15] ✓ 0.0s

... 335.5

```

The above code outputs an average # of epochs to convergence of 335.5 which is closest to 350, so the answer is a.

10)

An iteration of PLA is  $w_{new} = w_{old} + yx$  where  $(x,y)$  is a misclassified point. If the point is classified, then there is no change in  $w$ . With SGD, the form is

$w_{\text{new}} = w_{\text{old}} - \eta \nabla E_{\text{in}}(w)$ . Assuming the learning rate to be 1,

we need  $-\nabla E_{\text{in}}(w)$  to match  $yx$ . The gradient of answer choice  $e$  is  $-\min(0, y_n x_n)$  & plugging this into

$-\nabla E_{\text{in}}(w)$  gives us  $yx$  & the desired behavior for PLA, so

the answer is  $e$ .