

All the code can be accessed here and executed in google colab or any Jupyter Notebook environment: <https://github.com/DragonBoy25830/caltech-cs-156>

1) In general, H' will have a lower complexity than H , so H' won't be able to explain more of f and the deterministic noise will increase, so the answer is b.

2)

```
[1] import numpy as np
    ✓ 0.4s MagicPython

[2] training_set = []
    testing_set = []
    ✓ 0.0s MagicPython

[3] with open('data_in.txt', 'r') as f:
    for line in f:
        training_set.append([float(x) for x in line.split()])

    with open('data_out.txt', 'r') as f:
        for line in f:
            testing_set.append([float(x) for x in line.split()])
    ✓ 0.0s MagicPython

[4] np_training_set = np.array(training_set)
    np_testing_set = np.array(testing_set)
    ✓ 0.0s MagicPython

[5] def transform_data(x1, x2):
    return (1, x1, x2, x1**2, x2**2, x1 * x2, np.abs(x1 - x2), np.abs(x1 + x2))
    ✓ 0.0s MagicPython
```

```
def calc_error(classification, w, z_transform):
    yn = classification
    counter = 0
    for i in range(len(z_transform)):
        z = z_transform[i]
        w_val = np.sign(np.dot(w, z))

        if w_val != yn[i]:
            counter += 1

    return counter / len(classification)

[6] ✓ 0.0s MagicPython

def get_x_values(dataset):
    dataset = np.array(dataset)
    x1 = dataset[:, 0]
    x2 = dataset[:, 1]

    return np.column_stack((x1, x2))

[7] ✓ 0.0s MagicPython

def run_linear_regression_experiment():
    train_xn = get_x_values(training_set)
    train_yn = np_training_set[:, 2]

    test_xn = get_x_values(testing_set)
    test_yn = np_testing_set[:, 2]

    train_xn_transform = [transform_data(x1, x2) for (x1, x2) in train_xn]
    test_xn_transform = [transform_data(x1, x2) for (x1, x2) in test_xn]

    w = np.matmul(np.linalg.pinv(train_xn_transform), train_yn)

    E_in = calc_error(train_yn, w, train_xn_transform)
    E_out = calc_error(test_yn, w, test_xn_transform)

    return E_in, E_out

[8] ✓ 0.0s MagicPython
```

```
run_linear_regression_experiment()

[9] ✓ 0.0s

... (0.02857142857142857, 0.084)
```

The code above outputs an E_{in} of 0.0286 and E_{out} of 0.084 which is closest in Euclidean distance to 0.03 & 0.08, so the answer is a.

3) The functions from before are used for this part as well. All I did was change the code for calculating

as well. All I did was change the code for calculating the weights.

```
def calculate_w_reg(Z, ZT, lambda_value, y):
    step_1 = np.matmul(ZT, Z)
    step_2 = step_1 + 1 / (len(step_1))
    step_3 = np.linalg.pinv(step_2)
    step_4 = np.matmul(step_3, ZT)
    w_reg = np.matmul(step_4, y)
    return w_reg

def run_weight_decay_experiment(lambda_value):
    train_xn = get_x_values(training_set)
    train_yn = np_training_set[:, 2]

    test_xn = get_x_values(testing_set)
    test_yn = np_testing_set[:, 2]

    train_xn_transform = np.array([transform_data(x1, x2) for (x1, x2) in train_xn])
    test_xn_transform = np.array([transform_data(x1, x2) for (x1, x2) in test_xn])

    w_reg = calculate_w_reg(train_xn_transform, np.transpose(train_xn_transform), lambda_value)

    E_in = calc_error(train_yn, w_reg, train_xn_transform)
    E_out = calc_error(test_yn, w_reg, test_xn_transform)

    return E_in, E_out

run_weight_decay_experiment(1e-3)
```

... (0.02857142857142857, 0.08)

The code above outputs an E_{in} of 0.0286 and E_{out} of 0.08 which is closest in Euclidean distance to 0.03 & 0.08, so the answer is d.

4)

```
run_weight_decay_experiment(1e3)
```

[25] ✓ 0.0s

4)

```

run_weight_decay_experiment(1e3)
[25] ✓ 0.0s
... (0.37142857142857144, 0.436)

```

The code above outputs an E_{in} of 0.371 and E_{out} of 0.436 which is closest in Euclidean distance to 0.4 & 0.4, so the answer is e.

5)

```

k_values = [2, 1, 0, -1, -2]

for k in k_values:
    Ein, Eout = run_weight_decay_experiment(10 ** k)
    print(f"k = {k}: Eout = {Eout}")
[27] ✓ 0.0s
... k = 2: Eout = 0.228
    k = 1: Eout = 0.124
    k = 0: Eout = 0.092
    k = -1: Eout = 0.056
    k = -2: Eout = 0.084

```

The above code shows that the smallest $E_{out} = 0.056$ occurs when $k = -1$, so the answer is d.

6)

```

k_values = np.arange(-100, 100)
min_Eout = float('inf')

for k in k_values:
    Ein, Eout = run_weight_decay_experiment(np.power(10.0, k))
    if Eout < min_Eout:
        min_Eout = Eout

```

6)

```

k_values = np.arange(-100, 100)
min_Eout = float('inf')

for k in k_values:
    Ein, Eout = run_weight_decay_experiment(np.power(10.0, k))
    if Eout < min_Eout:
        min_Eout = Eout

min_Eout

```

[32] ✓ 0.2s
... 0.056

The code above calculates E_{out} for integer values of k from -100 to 100 . The minimum E_{out} is equal to 0.056 which is closest to 0.06 , so the answer is b

7) If $a < b$, then $\mathcal{H}_a \in \mathcal{H}_b$. Using this fact, we can write that based off the given constraint,

$\mathcal{H}(Q, C, Q_0) = \mathcal{H}_{Q_0-1}$ if $C=0$ since the constraint zeroes out all the terms of L_2 after for $Q > Q_0$. As such,

$\mathcal{H}(10, 0, 3) = \mathcal{H}_2$ & $\mathcal{H}(10, 0, 4) = \mathcal{H}_3$. Since $\mathcal{H}_2 \in \mathcal{H}_3$,

we can say that $\mathcal{H}(10, 0, 3) \cap \mathcal{H}(10, 0, 4) = \mathcal{H}_2$. As

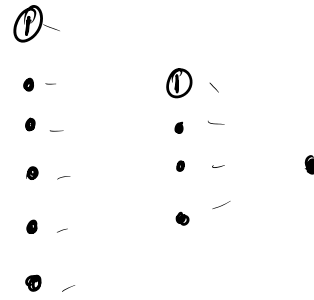
such, the answer is c.

8)

$$W_{ij}^{(k)} x_i^{(k-1)} = \text{calculate } x_j$$

$$W_{ij}^{(k)} \delta_j^{(k)} = \text{calculate } \delta_j$$

$$x_i^{(k-1)} \delta_j^{(k)} = \text{update weights}$$



We calculate the total # of operations by looking at how many times each type of operation is performed:

1) Calculate x_j

$$\sum_{i=0}^{d^{(L-1)}} w_{ij}^{(k)} x_i^{(k-1)} = \underbrace{\sum_{i=0}^{d^{(0)}} w_{ij}^{(1)} x_i^{(0)}}_{\text{6 weights for each of the 3 inputs in } d^{(0)}, \text{ so 18 ops.}} + \underbrace{\sum_{i=0}^{d^{(1)}} w_{ij}^{(2)} x_i^{(1)}}_{\text{4 weights for the 1 input in } d^{(1)}, \text{ so 4 ops.}}$$

6 weights for each of the 3 inputs in $d^{(0)}$, so 18 ops.

4 weights for the 1 input in $d^{(1)}$, so 4 ops.

2) Calculate δ_j - Backpropagation

$$\sum_{j=1}^{d^{(k)}} w_{ij}^{(k)} \delta_j^{(k)} = \underbrace{\sum_{j=1}^{d^{(1)}} w_{ij}^{(1)} \delta_j^{(1)}}_{\text{we only calculate } \delta_j \text{'s in the backward step for the four units/nodes in } d^{(1)} \text{ because } d^{(2)} \text{ is the base case}}$$

δ_j 's are only calculated for $d^{(1)}$, so there are 3

we only calculate δ_j 's in the backward step for the four units/nodes in $d^{(1)}$ because $d^{(2)}$ is the base case

so there are 3
S's calculated so
3 ops

$d^{(2)}$ is the base case

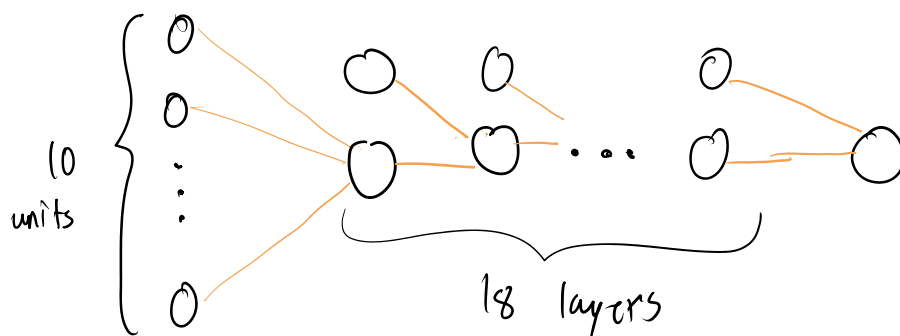
3) Update Weights

every w_{ij} is updated and there are 22
total weights, so there are 22 total ops.

$$\text{total \# of operations} = 22 + 3 + 22 = 47 \approx 45$$

The answer is d.

9) The simplest way to reduce the number of
weights is to only have two unit in each
hidden layer.



We have 10 weights from the input layer to the first hidden layer, 34 weights within the hidden layers and two weights from the last hidden layer to the output layer, so the minimum # of weights is $10+34+2=46$.

The answer is a.

10)

```
for i in range(1, 36):  
    print(i, 10 * (i - 1) + i * (36 - i - 1) + (36 - i))
```

[9] ✓ 0.0s

...	1	69
	2	110
	3	149
	4	186
	5	221
	6	254
	7	285
	8	314
	9	341
	10	366
	11	389
	12	410
	13	429
	14	446
	15	461
	16	474
	17	485
	18	494
	19	501
	20	506
	21	509
	22	510
	23	509
	24	506
	25	501
	...	
	32	410
	33	389
	34	366
	35	341

After some trial & error, it seems that two hidden layers maximizes the number of weights

in the network. After writing some python code, we see that 22 units in the first layer and 14 in the second layer results in a maximum # of weights of 510. The answer is e.

↓

$$10(21) + 22(13) + 14 = 510$$